

CS252
Graduate Computer Architecture
Lecture 2

**Review of Cost, Integrated Circuits, Benchmarks,
Moore's Law, & Prerequisite Quiz**

January 19, 2001
Prof. David A. Patterson
Computer Science 252
Spring 2001

Review #1/3: Pipelining & Performance

- Just overlap tasks; easy if tasks are independent
- Speed Up \leq Pipeline Depth; if ideal CPI is 1, then:

$$\text{Speedup} = \frac{\text{Pipeline depth}}{1 + \text{Pipeline stall CPI}} \times \frac{\text{Cycle Time}_{\text{unpipelined}}}{\text{Cycle Time}_{\text{pipelined}}}$$

- Hazards limit performance on computers:
 - Structural: need more HW resources
 - Data (RAW, WAR, WAW): need forwarding, compiler scheduling
 - Control: delayed branch, prediction
- Time is measure of performance: latency or throughput
- CPI Law:

CPU time	=	$\frac{\text{Seconds}}{\text{Program}}$	=	$\frac{\text{Instructions}}{\text{Program}}$	x	$\frac{\text{Cycles}}{\text{Instruction}}$	x	$\frac{\text{Seconds}}{\text{Cycle}}$
----------	---	---	---	--	---	--	---	---------------------------------------


Review #2/3: Caches

- The Principle of Locality:
 - Program access a relatively small portion of the address space at any instant of time.
 - » Temporal Locality: Locality in Time
 - » Spatial Locality: Locality in Space
- Three Major Categories of Cache Misses:
 - Compulsory Misses: sad facts of life. Example: cold start misses.
 - Capacity Misses: increase cache size
 - Conflict Misses: increase cache size and/or associativity.
- Write Policy:
 - Write Through: needs a write buffer.
 - Write Back: control can be complex
- Today CPU time is a function of (ops, cache misses) vs. just $f(\text{ops})$: What does this mean to Compilers, Data structures, Algorithms?

Now, Review of Virtual Memory


Basic Issues in VM System Design

- size of information blocks that are transferred from secondary to main storage (M)
- block of information brought into M, and M is full, then some region of M must be released to make room for the new block --> *replacement policy*
- which region of M is to hold the new block --> *placement policy*
- missing item fetched from secondary memory only on the occurrence of a fault --> *demand load policy*



Paging Organization

virtual and physical address space partitioned into blocks of equal size




Address Map

$V = \{0, 1, \dots, n - 1\}$ virtual address space $n > m$
 $M = \{0, 1, \dots, m - 1\}$ physical address space

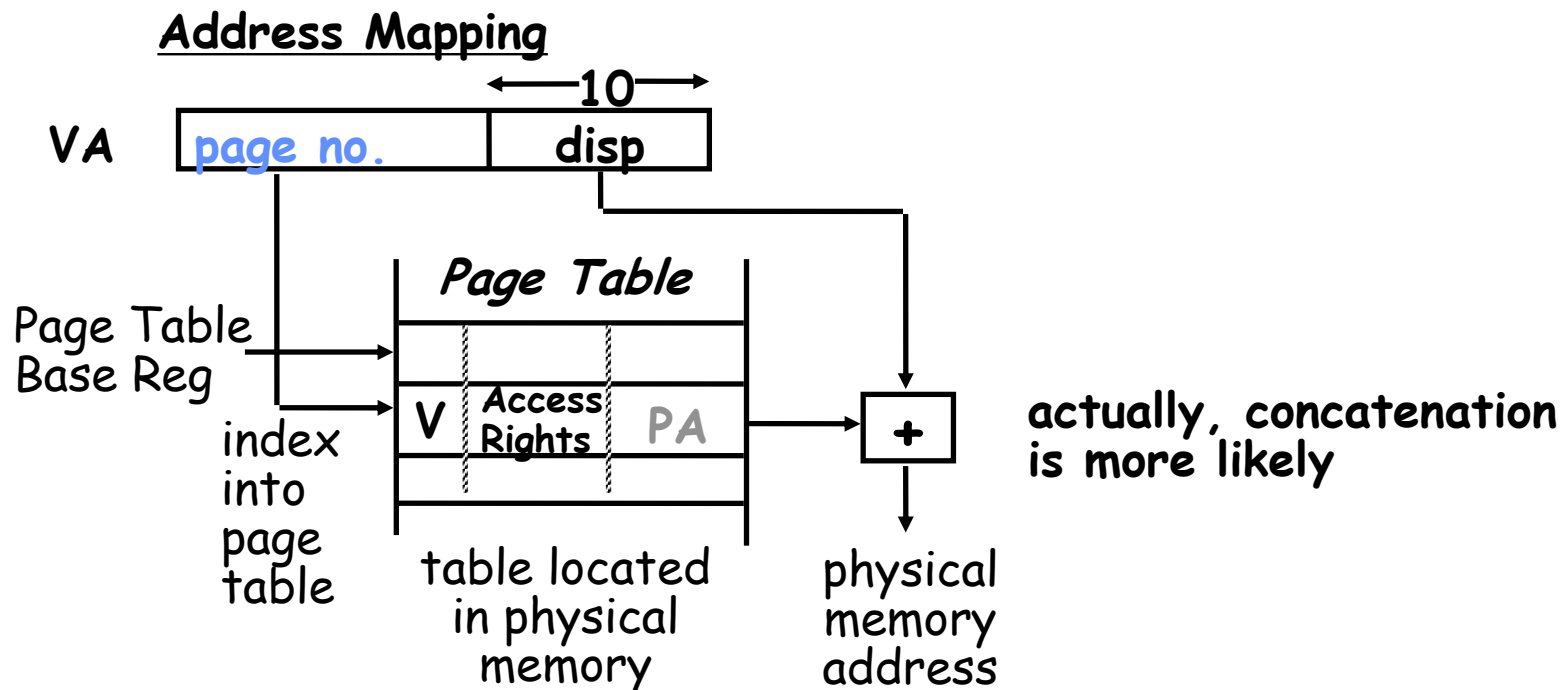

MAP: $V \dashrightarrow M \cup \{\emptyset\}$ address mapping function

$\text{MAP}(a) = a'$ if data at virtual address a is present in physical address a' and a' in M

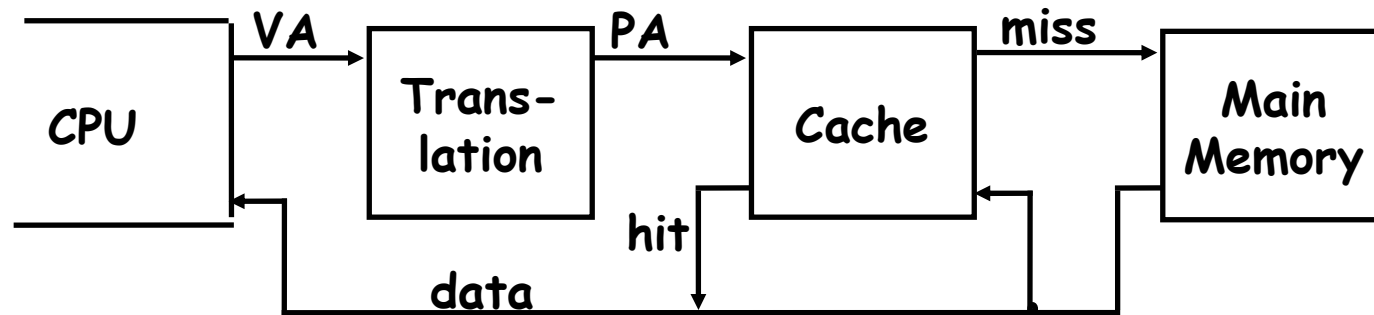
$= \emptyset$ if data at virtual address a is not present in M



Paging Organization



Virtual Address and a Cache



It takes an extra memory access to translate VA to PA

This makes cache access very expensive, and this is the "innermost loop" that you want to go as fast as possible

ASIDE: Why access cache with PA at all? VA caches have a problem!

synonym / alias problem: two different virtual addresses map to same physical address => two different cache entries holding data for the same physical address!

for update: must update all cache entries with same physical address or memory becomes inconsistent

determining this requires significant hardware: essentially an associative lookup on the physical address tags to see if you have multiple hits

or software enforced **alias boundary**: same lsb of VA & PA > cache size

TLBs

A way to speed up translation is to use a special cache of recently used page table entries -- this has many names, but the most frequently used is *Translation Lookaside Buffer* or *TLB*

Virtual Address	Physical Address	Dirty	Ref	Valid	Access


Really just a cache on the page table mappings

TLB access time comparable to cache access time
(much less than main memory access time)

Translation Look-Aside Buffers

Just like any other cache, the TLB can be organized as fully associative, set associative, or direct mapped

TLBs are usually small, typically not more than 128 - 256 entries even on high end machines. This permits fully associative lookup on these machines. Most mid-range machines use small n-way set associative organizations.



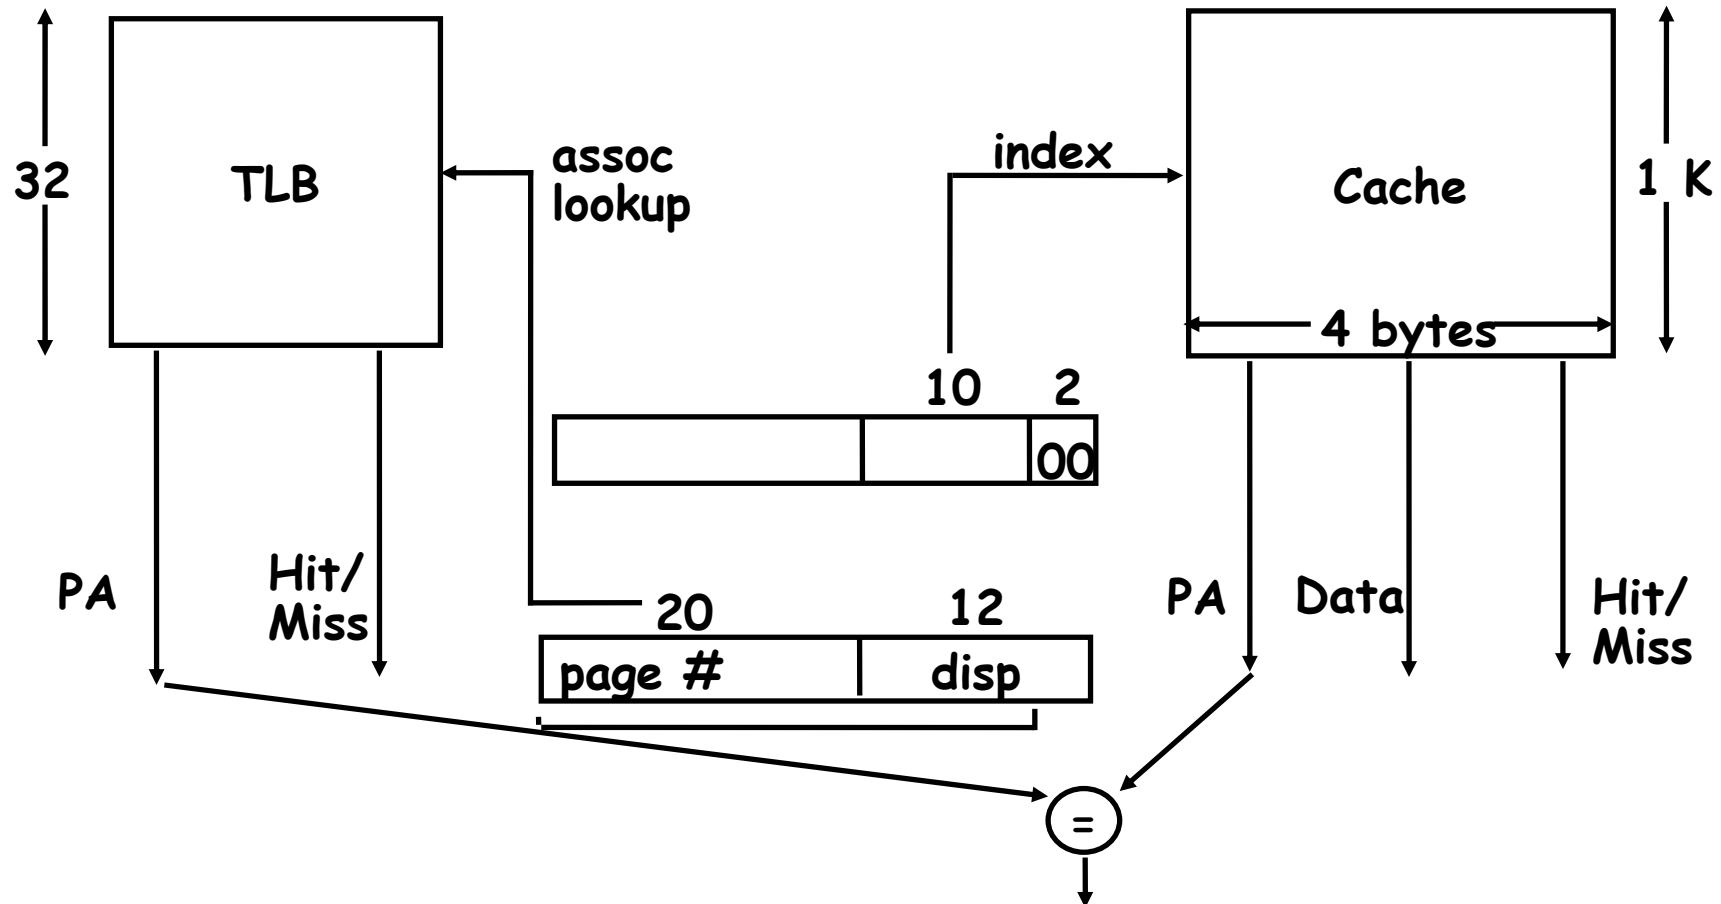
Reducing Translation Time

Machines with TLBs go one step further to reduce # cycles/cache access

They overlap the cache access with the TLB access:

high order bits of the VA are used to look in the TLB while low order bits are used as index into cache

Overlapped Cache & TLB Access



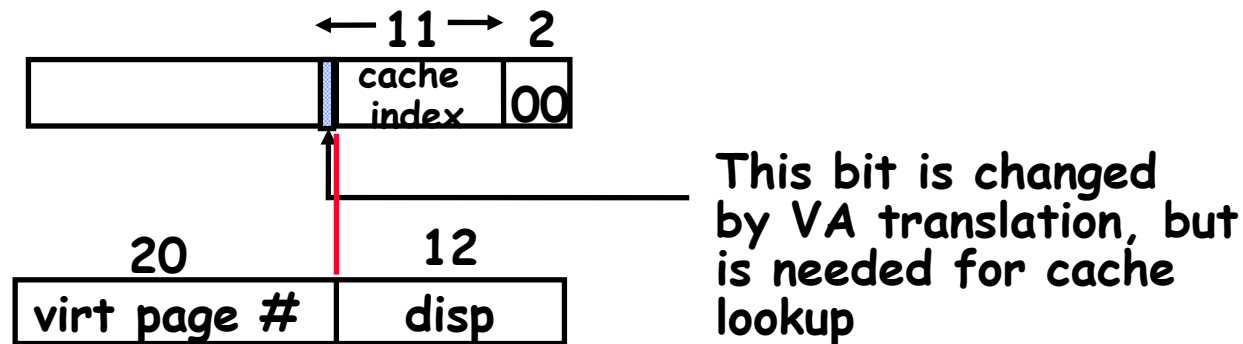
IF cache hit AND (cache tag = PA) then deliver data to CPU
 ELSE IF [cache miss OR (cache tag = PA)] and TLB hit THEN
 access memory with the PA from the TLB
 ELSE do standard VA translation

Problems With Overlapped TLB Access

Overlapped access only works as long as the address bits used to index into the cache *do not change* as the result of VA translation


This usually limits things to small caches, large page sizes, or high n-way set associative caches if you want a large cache

Example: suppose everything the same except that the cache is increased to 8 K bytes instead of 4 K:



Solutions:

go to 8K byte page sizes;
go to 2 way set associative cache; or
SW guarantee $VA[13]=PA[13]$



SPEC: System Performance Evaluation Cooperative

- **First Round 1989**
 - 10 programs yielding a single number ("SPECmarks")
- **Second Round 1992**
 - SPECInt92 (6 integer programs) and SPECfp92 (14 floating point programs)
 - » Compiler Flags unlimited. March 93 of DEC 4000 Model 610:
`spice: unix.c:/def=(sysv,has_bcopy,"bcopy(a,b,c)=memcpy(b,a,c)"`
`wave5: /ali=(all,dcom=nat)/ag=a/ur=4/ur=200`
`nasa7: /norecu/ag=a/ur=4/ur2=200/lc=blas`
- **Third Round 1995**
 - new set of programs: SPECint95 (8 integer programs) and SPECfp95 (10 floating point)
 - "benchmarks useful for 3 years"
 - Single flag setting for all programs: SPECint_base95, SPECfp_base95

SPEC: System Performance Evaluation Cooperative

- Fourth Round 2000: SPEC CPU2000
 - 12 Integer
 - 14 Floating Point
 - 2 choices on compilation; “aggressive” (SPECint2000, SPECfp2000), “conservative” (SPECint_base2000, SPECfp_base); flags same for all programs, no more than 4 flags, same compiler for conservative, can change for aggressive
 - multiple data sets so that can train compiler if trying to collect data for input to compiler to improve optimization

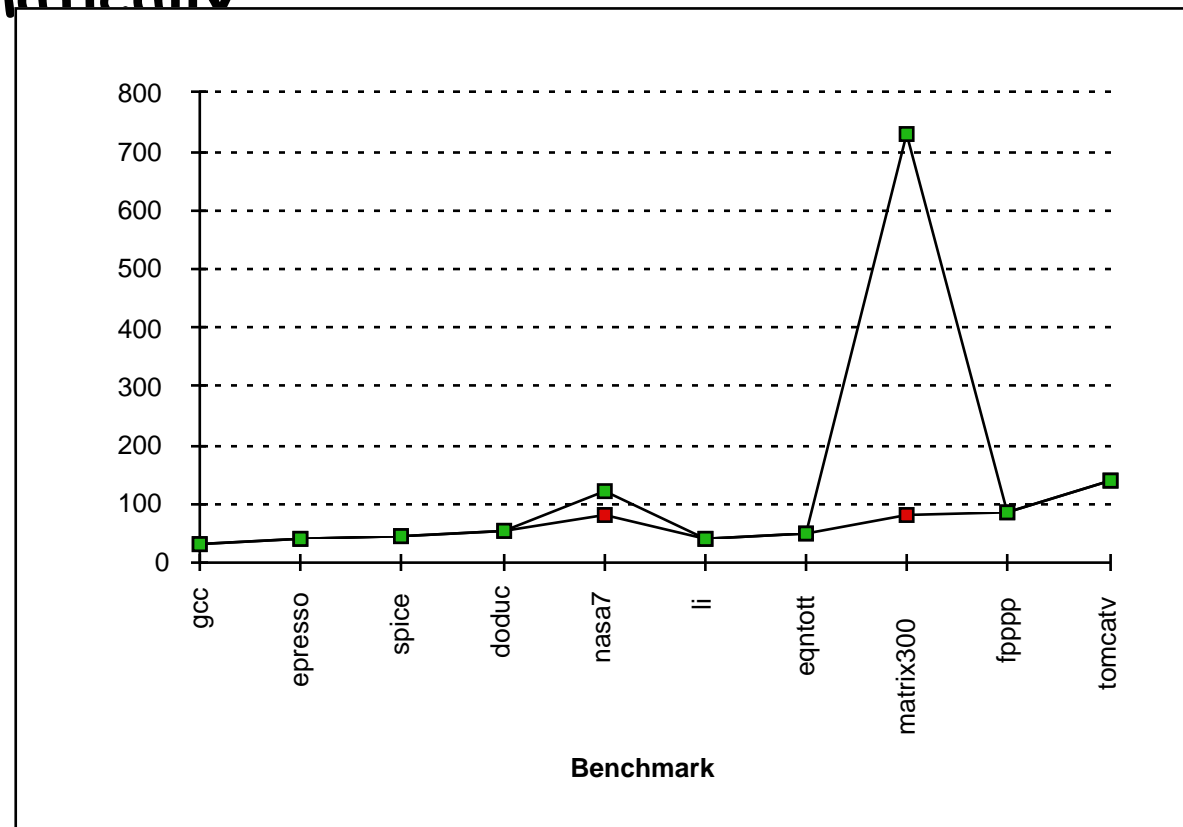
How to Summarize Performance

- Arithmetic mean (weighted arithmetic mean) tracks execution time:
 $\Sigma(T_i)/n$ or $\Sigma(W_i * T_i)$
- Harmonic mean (weighted harmonic mean) of rates (e.g., MFLOPS) tracks execution time:
 $n/\Sigma(1/R_i)$ or $n/\Sigma(W_i/R_i)$
- Normalized execution time is handy for scaling performance (e.g., X times faster than SPARCstation 10)
- But do not take the arithmetic mean of normalized execution time, use the geometric mean:

$$(\prod T_j / N_j)^{1/n}$$

SPEC First Round

- One program: 99% of time in single line of code
- New front-end compiler could improve dramatically



Impact of Means on SPECmark89 for IBM 550

	Ratio to VAX:		<u>Time:</u>		<u>Weighted Time:</u>	
<i>Program</i>	<i>Before</i>	<i>After</i>	<i>Before</i>	<i>After</i>	<i>Before</i>	<i>After</i>
gcc	30	29	49	51	8.91	9.22
espresso	35	34	65	67	7.64	7.86
spice	47	47	510	510	5.69	5.69
doduc	46	49	41	38	5.81	5.45
nasa7	78	144	258	140	3.43	1.86
li	34	34	183	183	7.86	7.86
eqntott	40	40	28	28	6.68	6.68
matrix300	78	730	58	6	3.43	0.37
fpppp	90	87	34	35	2.97	3.07
tomcatv	33	138	20	19	2.01	1.94
Mean	54	72	124	108	54.42	49.99
	Geometric		Arithmetic		Weighted Arith.	
	Ratio	1.33	Ratio	1.16	Ratio	1.09

Performance Evaluation


- “For better or worse, benchmarks shape a field”
- Good products created when have:
 - Good benchmarks
 - Good ways to summarize performance
- Given sales is a function in part of performance relative to competition, investment in improving product as reported by performance summary
- If benchmarks/summary inadequate, then choose between improving product for real programs vs. improving product to get more sales; Sales almost always wins!
- Execution time is the measure of computer performance!

Integrated Circuits Costs

$$\text{IC cost} = \frac{\text{Die cost} + \text{Testing cost} + \text{Packaging cost}}{\text{Final test yield}}$$

$$\text{Die cost} = \frac{\text{Wafer cost}}{\text{Dies per Wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi (\text{Wafer_diam}/2)^2}{\text{Die_Area}} - \frac{\pi \times \text{Wafer_diam}}{\sqrt{2} \cdot \text{Die_Area}} - \text{Test_Die}$$



$$\text{Die Yield} = \text{Wafer_yield} \times \left\{ 1 + \left(\frac{\text{Defect_Density} \times \text{Die_area}}{\alpha} \right)^{-\alpha} \right\}$$

Die Cost goes roughly with die area⁴

Real World Examples

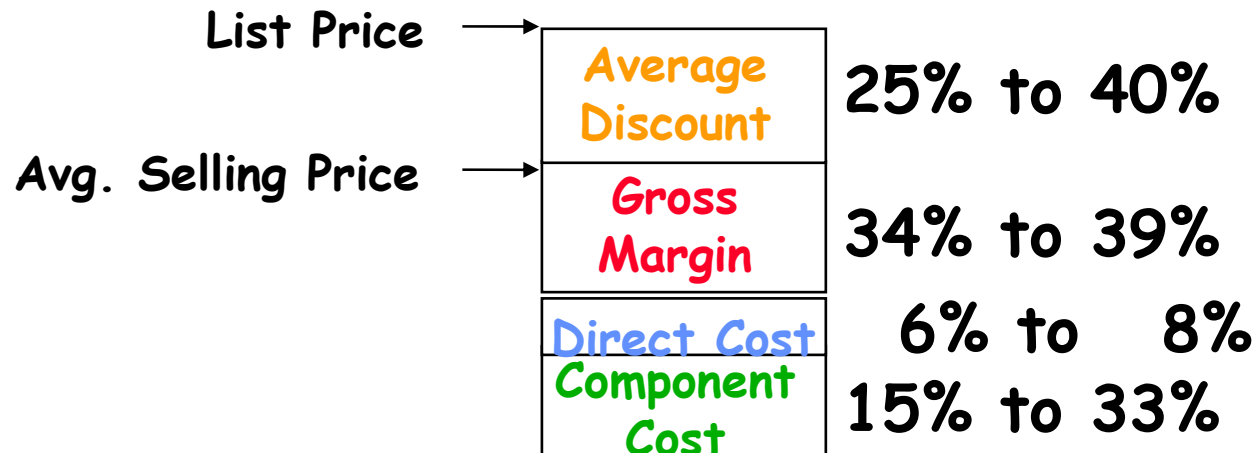
Chip	Metal layers	Line width	Wafer cost	Defect /cm ²	Area mm ²	Dies/ wafer	Yield	Die Cost
386DX	2	0.90	\$900	1.0	43	360	71%	\$4
486DX2	3	0.80	\$1200	1.0	81	181	54%	\$12
PowerPC 601	4	0.80	\$1700	1.3	121	115	28%	\$53
HP PA 7100	3	0.80	\$1300	1.0	196	66	27%	\$73
DEC Alpha	3	0.70	\$1500	1.2	234	53	19%	\$149
SuperSPARC	3	0.70	\$1700	1.6	256	48	13%	\$272
Pentium	3	0.80	\$1500	1.5	296	40	9%	\$417

– From "Estimating IC Manufacturing Costs," by Linley Gwennap, *Microprocessor Report*, August 2, 1993, p. 15

Cost/Performance

What is Relationship of Cost to Price?

- **Component Costs**
- **Direct Costs** (add 25% to 40%) recurring costs: labor, purchasing, scrap, warranty
- **Gross Margin** (add 82% to 186%) nonrecurring costs: R&D, marketing, sales, equipment maintenance, rental, financing cost, pretax profits, taxes
- **Average Discount** to get List Price (add 33% to 66%): volume discounts and/or retailer markup





Chip Prices (August 1993)

- Assume purchase 10,000 units

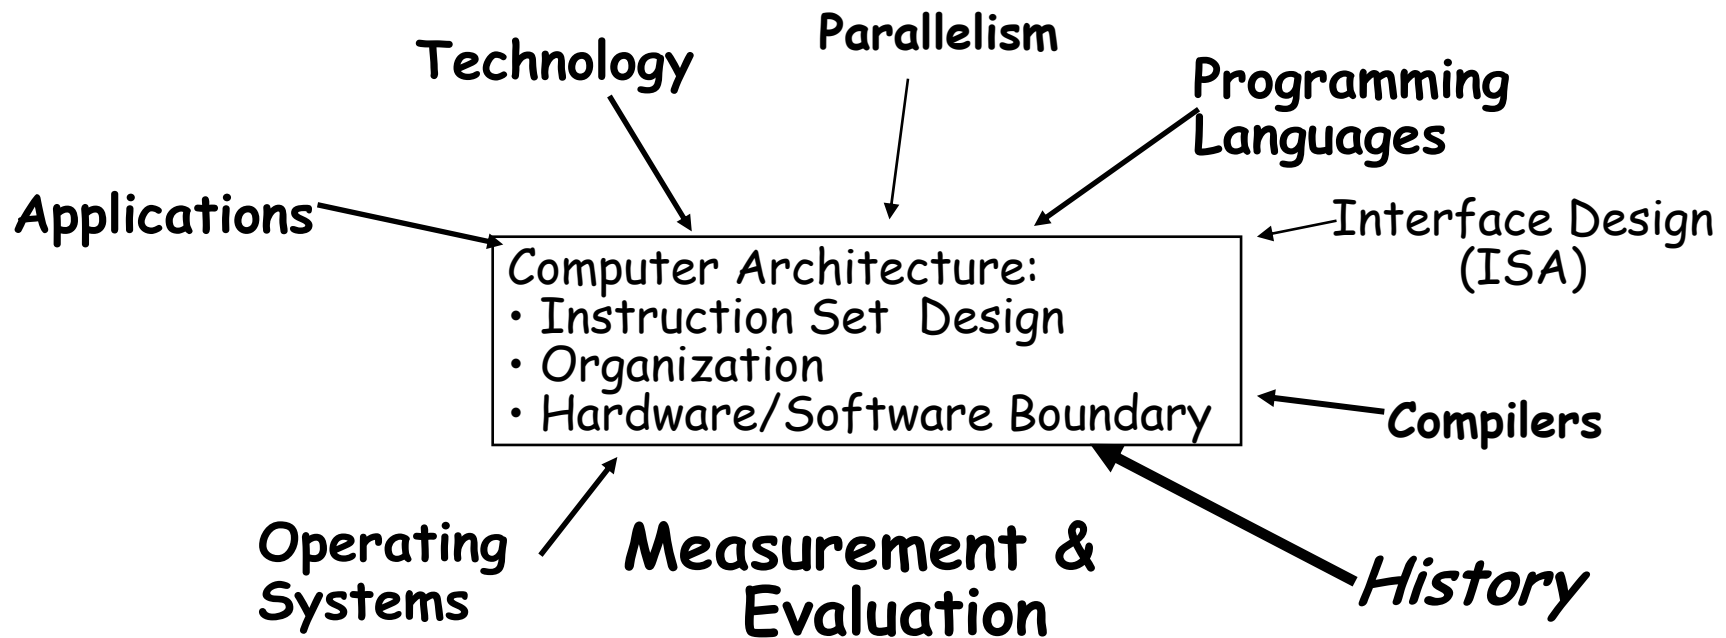
Chip	Area mm ²	Mfg. cost	Price	Multi- plier	Comment
386DX	43	\$9	\$31	3.4	Intense Competition
486DX2	81	\$35	\$245	7.0	No Competition
PowerPC 601	121	\$77	\$280	3.6	
DEC Alpha	234	\$202	\$1231	6.1	Recoup R&D?
Pentium	296	\$473	\$965	2.0	Early in shipments

Summary: Price vs. Cost



CS 252 Course Focus

Understanding the design techniques, machine structures, technology factors, evaluation methods that will determine the form of computers in 21st Century



Topic Coverage

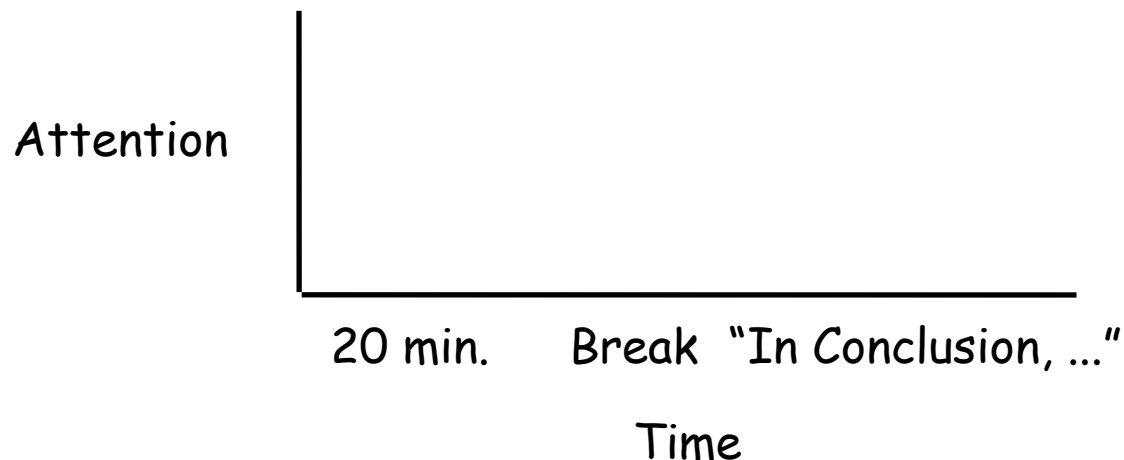
Textbook: Hennessy and Patterson, *Computer Architecture: A Quantitative Approach*, 3rd Ed., 2001

Research Papers -- Handed out in class

- 1 week: **Review**: Fundamentals of Computer Architecture (Ch. 1), Pipelining, Performance, Caches, Virtual Memory, Cost, Ics
- 1 week: Memory Hierarchy (Chapter 5)
- 2 weeks: Fault Tolerance, Queuing Theory, Input/Output and Storage (Ch. 6)
- 2 weeks: Networks and Clusters (Ch. 7)
- 2 weeks: Multiprocessors (Ch. 8)
- 2 weeks: Instruction Sets, DSPs, SIMD (Ch. 2), Vector Processors (Appendix B).
- 1 week: Dynamic Execution. (Ch 3)
- 1 week: Static Execution. (Ch 4)
- Rest: Project strategy meetings, presentations, quizzes

Lecture style


- 1-Minute Review
- 20-Minute Lecture/Discussion
- 5- Minute Administrative Matters
- 25-Minute Lecture/Discussion
- 5- Minute Class Discussion or Break (water, stretch)
- 25-Minute Lecture/Discussion
- Instructor will come to class early & stay after to answer questions



Quizzes


- Reduce the pressure of taking quizzes
 - Only 2 Graded Quizzes:
Tentative: Wed Mar 7th and Wed. Apr 18th
 - Our goal: test knowledge vs. speed writing
 - 3 hrs to take 1.5-hr test (5:30-8:30 PM, TBA location)
 - Both mid-term quizzes can bring summary sheet
 - » Transfer ideas from book to paper
 - Last chance Q&A: during class time day of exam
- Students/Faculty meet over free pizza/drinks at La Vals:
Wed Oct. 18th (8:30 PM) and Wed Apr 18th (8:30 PM)


Original




Big Fishes Eating Little Fishes

1988 Computer Food Chain





Massively Parallel Processors




Mini-
supercomputer




Mini-
computer


1998 Computer Food Chain



Mainframe



Supercomputer



Server




Work-
station PC

Now who is eating whom?

Why Such Change in 10 years?

- Performance
 - Technology Advances
 - » CMOS VLSI dominates older technologies (TTL, ECL) in cost **AND** performance
 - Computer architecture advances improves low-end
 - » RISC, superscalar, RAID, ...
- Price: Lower costs due to ...
 - Simpler development
 - » CMOS VLSI: smaller systems, fewer components
 - Higher volumes
 - » CMOS VLSI : same dev. cost 10,000 vs. 10,000,000 units
 - Lower margins by class of computer, due to fewer services
- Function
 - Rise of networking/local interconnection technology

Technology Trends: Microprocessor Capacity




"Graduation Window"

Alpha 21264: 15 million
Pentium Pro: 5.5 million
PowerPC 620: 6.9 million
Alpha 21164: 9.3 million
Sparc Ultra: 5.2 million

CMOS improvements:

- Die size: 2X every 3 yrs
- Line width: halve / 7 yrs

Memory Capacity (Single Chip DRAM)




year	size(Mb)	cyc time
1980	0.0625	250 ns
1983	0.25	220 ns
1986	1	190 ns
1989	4	165 ns
1992	16	145 ns
1996	64	120 ns
2000	256	100 ns


Technology Trends (Summary)

	Capacity	Speed (latency)
Logic	2x in 3 years	2x in 3 years
DRAM	4x in 3-4 years	2x in 10 years
Disk	4x in 2-3 years	2x in 10 years

Processor Performance Trends



Processor Performance (1.35X before, 1.55X now)



Performance Trends (Summary)

- Workstation performance (measured in Spec Marks) improves roughly 50% per year (2X every 18 months)
- Improvement in cost performance estimated at 70% per year

Moore's Law Paper

- Discussion
- What did Moore predict?
- 35 years later, how did it hold up?
- In your view, what was biggest surprise in paper?

Review #3/3: TLB, Virtual Memory

- Caches, TLBs, Virtual Memory all understood by examining how they deal with 4 questions: 1) Where can block be placed? 2) How is block found? 3) What block is replaced on miss? 4) How are writes handled?
- Page tables map virtual address to physical address
- TLBs make virtual memory practical
 - Locality in data => locality in addresses of data, temporal and spatial
- TLB misses are significant in processor performance
 - funny times, as most systems can't access all of 2nd level cache without TLB misses!
- Today VM allows many processes to share single memory without having to swap all processes to disk; today VM protection is more important than memory hierarchy

Summary

- Performance Summary needs good benchmarks and good ways to summarize performance
- Transistors/chip for microprocessors growing via "Moore's Law" 2X 1.5/yr
- Disk capacity (so far) is at a faster rate last 4-5 years
- DRAM capacity is at a slower rate last 4-5 years
- In general, Bandwidth improving fast, latency improving slowly