

CS252
Graduate Computer Architecture
Lecture 4


Caches and Memory Systems

January 26, 2001
Prof. John Kubiatowicz

Review: Who Cares About the Memory Hierarchy?

- Processor Only Thus Far in Course:

- CPU cost/performance, ISA, Pipelined Execution



- 1980: no cache in μ proc; 1995 2-level cache on chip (1989 first Intel μ proc with a cache on chip)

Review: Cache performance

- Miss-oriented Approach to Memory Access:

$$CPUtime = IC \times \left(CPI_{Execution} + \frac{MemAccess}{Inst} \times MissRate \times MissPenalty \right) \times CycleTime$$

$$CPUtime = IC \times \left(CPI_{Execution} + \frac{MemMisses}{Inst} \times MissPenalty \right) \times CycleTime$$

- $CPI_{Execution}$ includes ALU and Memory instructions

- Separating out Memory component entirely

- AMAT = Average Memory Access Time

- CPI_{ALUOps} does not include memory instructions

$$CPUtime = IC \times \left(\frac{AluOps}{Inst} \times CPI_{AluOps} + \frac{MemAccess}{Inst} \times AMAT \right) \times CycleTime$$

$$\begin{aligned} AMAT &= HitTime + MissRate \times MissPenalty \\ &= (HitTime_{Inst} + MissRate_{Inst} \times MissPenalty_{Inst}) + \\ &\quad (HitTime_{Data} + MissRate_{Data} \times MissPenalty_{Data}) \end{aligned}$$

Review: Reducing Misses

- **Classifying Misses: 3 Cs**
 - **Compulsory**—Misses in even an Infinite Cache
 - **Capacity**—Misses in Fully Associative Size X Cache
 - **Conflict**—Misses in N-way Associative, Size X Cache
- **More recent, 4th "C":**
 - **Coherence** - Misses caused by cache coherence.

Review: Miss Rate Reduction

$$AMAT = HitTime + \text{MissRate} \times MissPenalty$$

- **3 Cs: Compulsory, Capacity, Conflict**
 1. Reduce Misses via Larger Block Size
 2. Reduce Misses via Higher Associativity
 3. Reducing Misses via Victim Cache
 4. Reducing Misses via Pseudo-Associativity
 5. Reducing Misses by HW Prefetching Instr, Data
 6. Reducing Misses by SW Prefetching Data
 7. Reducing Misses by Compiler Optimizations
- **Prefetching comes in two flavors:**
 - Binding prefetch: Requests load directly into register.
 - » Must be correct address and register!
 - Non-Binding prefetch: Load into cache.
 - » Can be incorrect. Frees HW/SW to guess!

Improving Cache Performance Continued

1. Reduce the miss rate,
2. Reduce the miss penalty, or
3. Reduce the time to hit in the cache.

$$AMAT = HitTime + MissRate \times MissPenalty$$

What happens on a Cache miss?

- For in-order pipeline, 2 options:

- Freeze pipeline in Mem stage (popular early on: Sparc, R4000)

IF	ID	EX	Mem	stall	stall	stall	...	stall	Mem	Wr
	IF	ID	EX	stall	stall	stall	...	stall	stall	Ex Wr

- Use Full/Empty bits in registers + MSHR queue

- » MSHR = "Miss Status/Handler Registers" (Kroft)
Each entry in this queue keeps track of status of outstanding memory requests to one complete memory line.

- Per cache-line: keep info about memory address.
- For each word: register (if any) that is waiting for result.
- Used to "merge" multiple requests to one memory line

- » New load creates MSHR entry and sets destination register to "Empty". Load is "released" from pipeline.

- » Attempt to use register before result returns causes instruction to block in decode stage.

- » Limited "out-of-order" execution with respect to loads.

Popular with in-order superscalar architectures.

- Out-of-order pipelines already have this functionality built in... (load queues, etc).

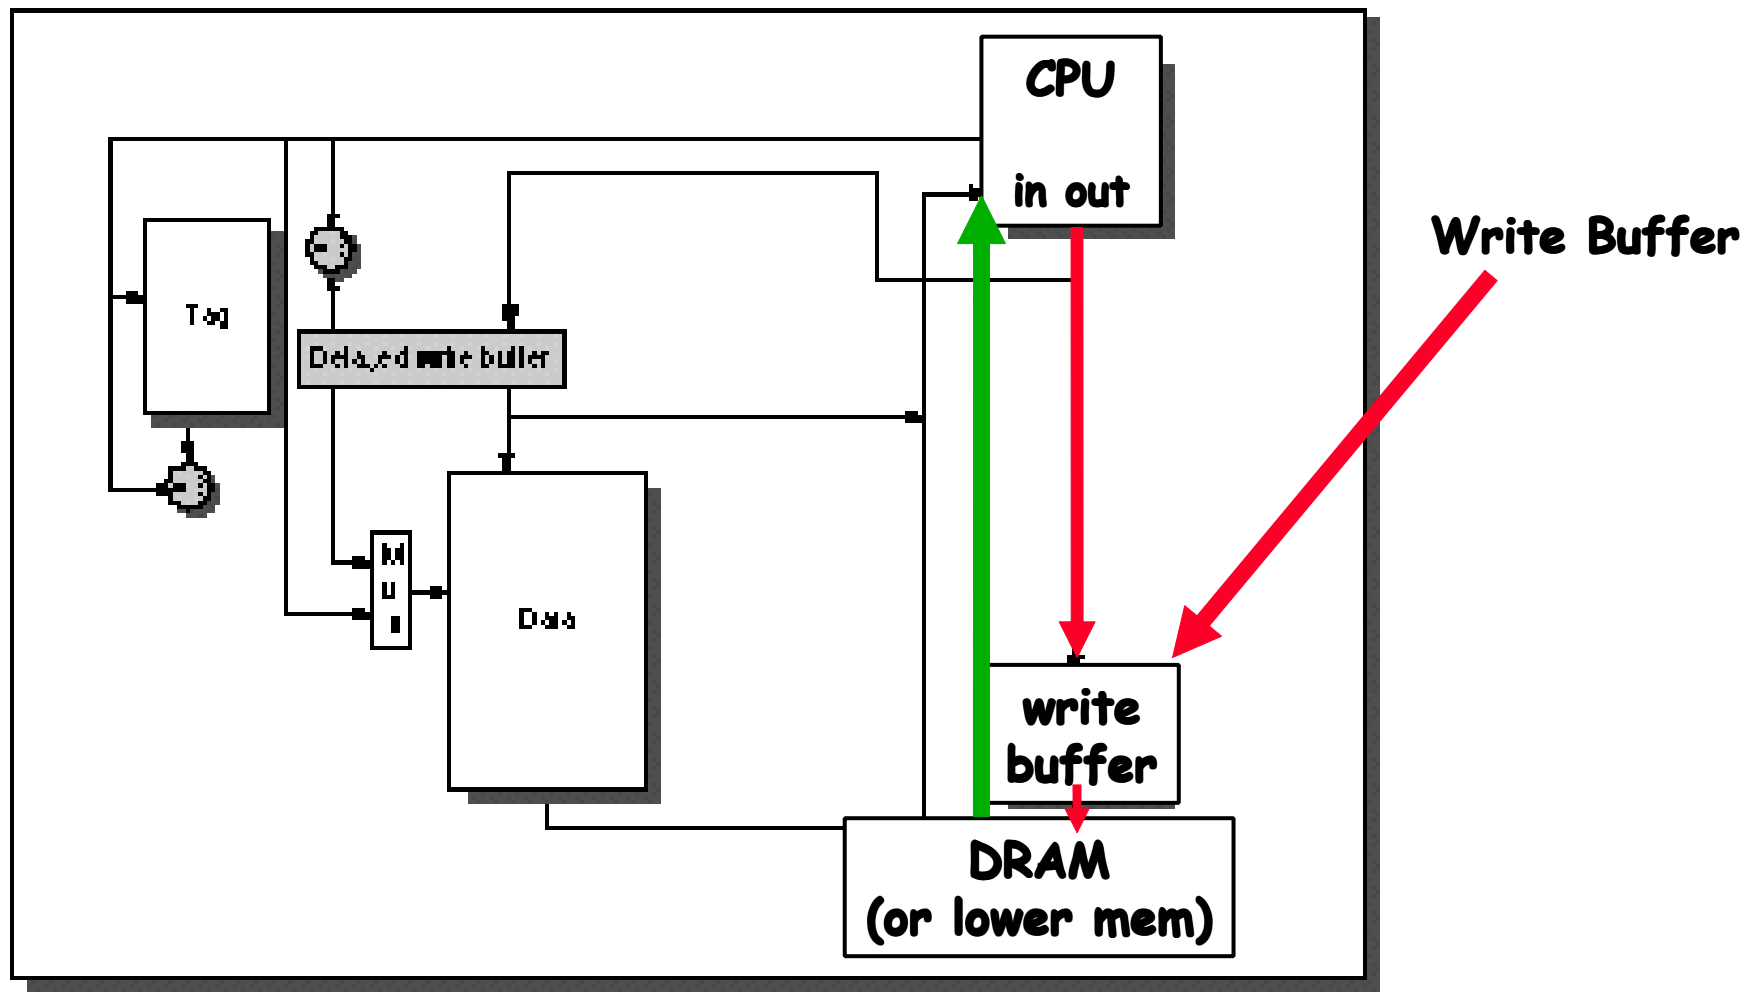
Write Policy: Write-Through vs Write-Back

- **Write-through:** all writes update cache and underlying memory/cache
 - Can always discard cached data - most up-to-date data is in memory
 - Cache control bit: only a *valid* bit
- **Write-back:** all writes simply update cache
 - Can't just discard cached data - may have to write it back to memory
 - Cache control bits: both *valid* and *dirty* bits
- **Other Advantages:**
 - **Write-through:**
 - » memory (or other processors) always have latest data
 - » Simpler management of cache
 - **Write-back:**
 - » much lower bandwidth, since data often overwritten multiple times
 - » Better tolerance to long-latency memory?

Write Policy 2: Write Allocate vs Non-Allocate (What happens on write-miss)

- **Write allocate: allocate new cache line in cache**
 - Usually means that you have to do a “read miss” to fill in rest of the cache-line!
 - Alternative: per/word valid bits
- **Write non-allocate (or “write-around”):**
 - Simply send write data through to underlying memory/cache - don't allocate new cache line!

1. Reducing Miss Penalty: Read Priority over Write on Miss



1. Reducing Miss Penalty: Read Priority over Write on Miss

- **Write-through with write buffers offer RAW conflicts with main memory reads on cache misses**
 - If simply wait for write buffer to empty, might increase read miss penalty (old MIPS 1000 by 50%)
 - Check write buffer contents before read; if no conflicts, let the memory access continue
- **Write-back also want buffer to hold misplaced blocks**
 - Read miss replacing dirty block
 - Normal: Write dirty block to memory, and then do the read
 - Instead copy the dirty block to a write buffer, then do the read, and then do the write
 - CPU stall less since restarts as soon as do read

2. Reduce Miss Penalty: Early Restart and Critical Word First

- Don't wait for full block to be loaded before restarting CPU
 - Early restart—As soon as the requested word of the block arrives, send it to the CPU and let the CPU continue execution
 - Critical Word First—Request the missed word first from memory and send it to the CPU as soon as it arrives; let the CPU continue execution while filling the rest of the words in the block. Also called *wrapped fetch* and *requested word first*
- Generally useful only in large blocks,
- Spatial locality a problem; tend to want next sequential word, so not clear if benefit by early restart




block

3. Reduce Miss Penalty: Non-blocking Caches to reduce stalls on misses

- Non-blocking cache or lockup-free cache allow data cache to continue to supply cache hits during a miss
 - requires F/E bits on registers or out-of-order execution
 - requires multi-bank memories
- “hit under miss” reduces the effective miss penalty by working during miss vs. ignoring CPU requests
- “hit under multiple miss” or “miss under miss” may further lower the effective miss penalty by overlapping multiple misses
 - Significantly increases the complexity of the cache controller as there can be multiple outstanding memory accesses
 - Requires **multiple memory banks** (otherwise cannot support)
 - Pentium Pro allows 4 outstanding memory misses

Value of Hit Under Miss for SPEC (Normalized to blocking cache)

Hit Under i Misses



- FP programs on average: **AMAT= 0.68 -> 0.52 -> 0.34 -> 0.26**
- Int programs on average: **AMAT= 0.24 -> 0.20 -> 0.19 -> 0.19**
- **8 KB Data Cache, Direct Mapped, 32B block, 16 cycle miss**

4. Second level cache

- L2 Equations

$$AMAT = \text{Hit Time}_{L1} + \text{Miss Rate}_{L1} \times \text{Miss Penalty}_{L1}$$

$$\text{Miss Penalty}_{L1} = \text{Hit Time}_{L2} + \text{Miss Rate}_{L2} \times \text{Miss Penalty}_{L2}$$

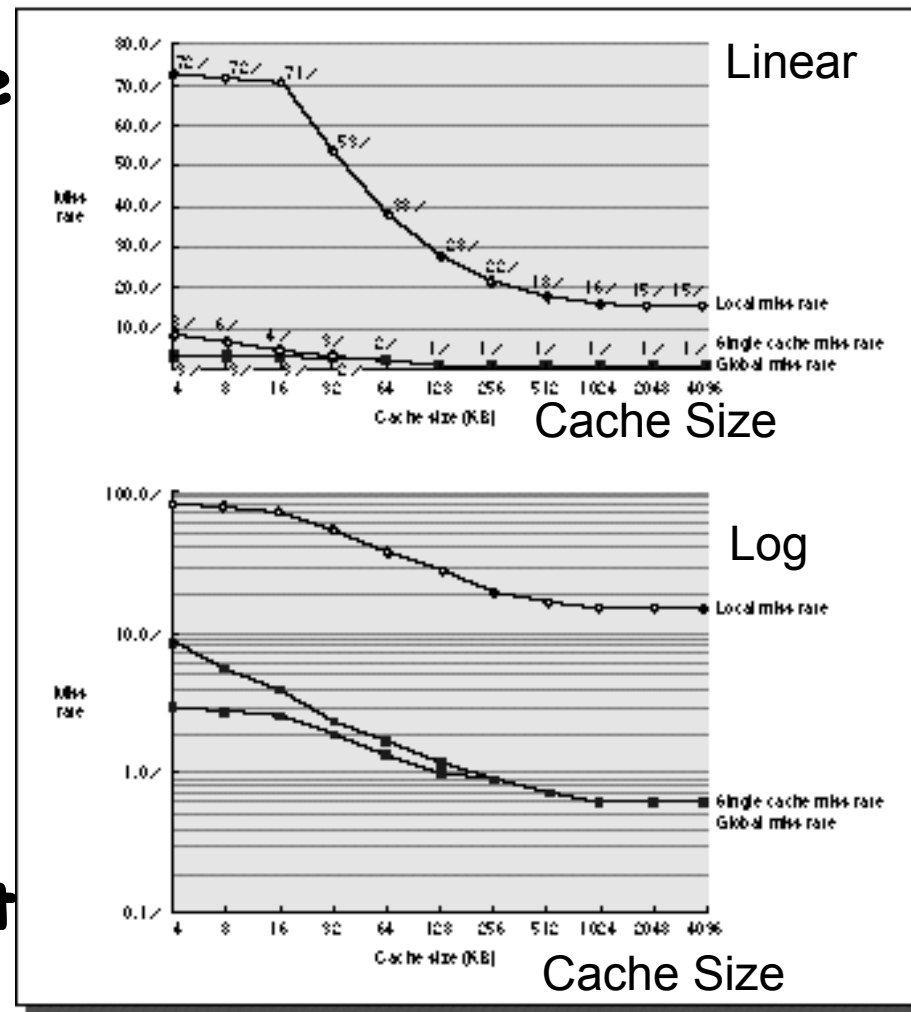
$$AMAT = \text{Hit Time}_{L1} + \text{Miss Rate}_{L1} \times (\text{Hit Time}_{L2} + \text{Miss Rate}_{L2} \times \text{Miss Penalty}_{L2})$$

- Definitions:

- **Local miss rate**— misses in this cache divided by the total number of memory accesses **to this cache** (Miss rate_{L2})
- **Global miss rate**—misses in this cache divided by the total number of memory accesses **generated by the CPU** ($\text{Miss Rate}_{L1} \times \text{Miss Rate}_{L2}$)
- Global Miss Rate is what matters

Comparing Local and Global Miss Rates

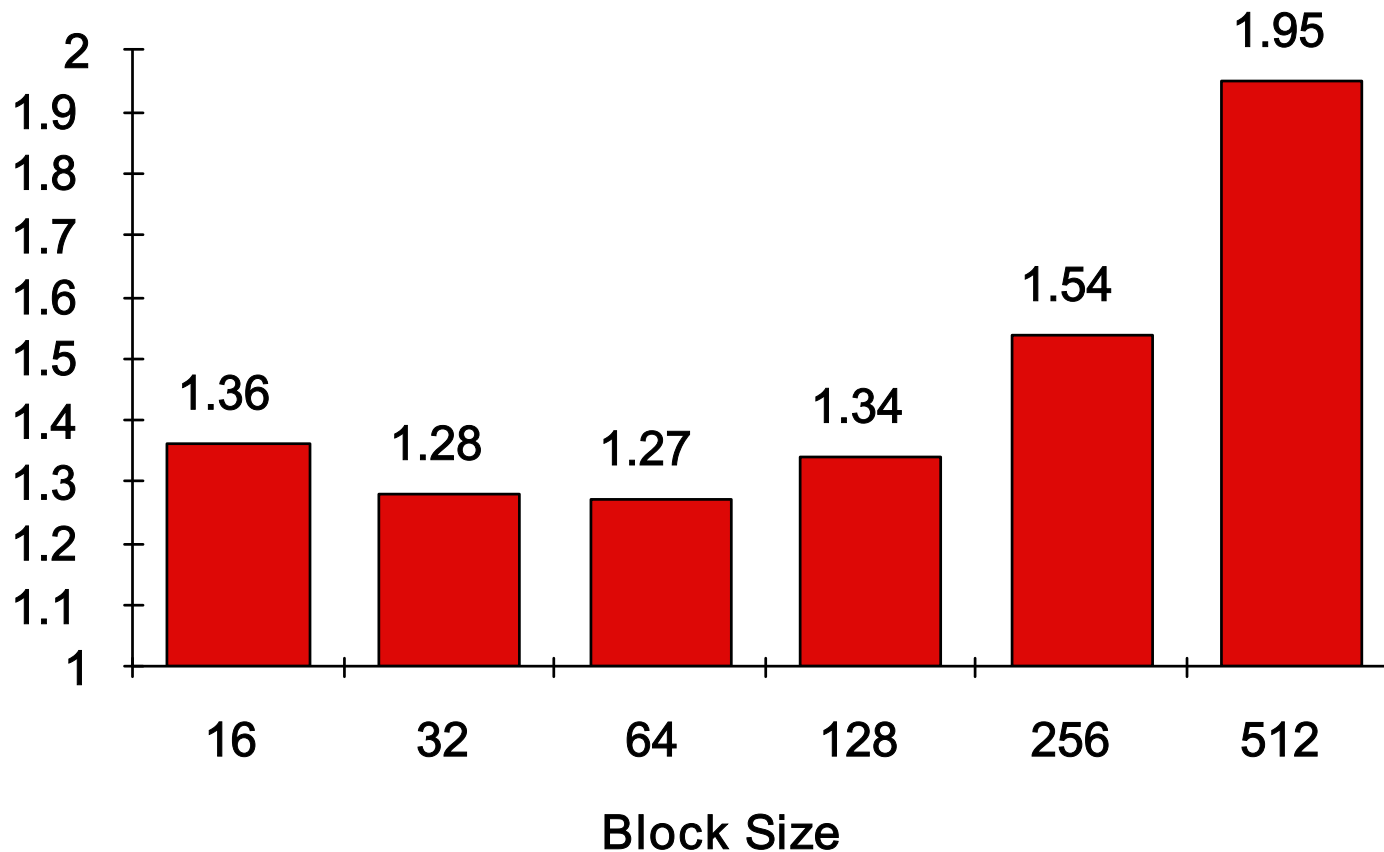
- 32 KByte 1st level cache; Increasing 2nd level cache
- Global miss rate close to single level cache rate provided $L2 \gg L1$
- Don't use local miss rate
- L2 not tied to CPU clock cycle!
- Cost & A.M.A.T.
- Generally Fast Hit Times and fewer misses
- Since hits are few, target miss reduction



Reducing Misses: Which apply to L2 Cache?

- **Reducing Miss Rate**
 1. **Reduce Misses via Larger Block Size**
 2. **Reduce Conflict Misses via Higher Associativity**
 3. **Reducing Conflict Misses via Victim Cache**
 4. **Reducing Conflict Misses via Pseudo-Associativity**
 5. **Reducing Misses by HW Prefetching Instr, Data**
 6. **Reducing Misses by SW Prefetching Data**
 7. **Reducing Capacity/Conf. Misses by Compiler Optimizations**

L2 cache block size & A.M.A.T. Relative CPU Time



- 32KB L1, 8 byte path to memory

Reducing Miss Penalty Summary

$$AMAT = HitTime + MissRate \times MissPenalty$$


- **Four techniques**
 - Read priority over write on miss
 - Early Restart and Critical Word First on miss
 - Non-blocking Caches (Hit under Miss, Miss under Miss)
 - Second Level Cache
- **Can be applied recursively to Multilevel Caches**
 - Danger is that time to DRAM will grow with multiple levels in between
 - First attempts at L2 caches can make things worse, since increased worst case is worse

Main Memory Background


- Performance of Main Memory:
 - **Latency**: Cache Miss Penalty
 - » **Access Time**: time between request and word arrives
 - » **Cycle Time**: time between requests
 - **Bandwidth**: I/O & Large Block Miss Penalty (L2)
- Main Memory is **DRAM**: Dynamic Random Access Memory
 - Dynamic since needs to be **refreshed** periodically (8 ms, 1% time)
 - Addresses divided into 2 halves (Memory as a 2D matrix):
 - » **RAS** or **Row Access Strobe**
 - » **CAS** or **Column Access Strobe**
- Cache uses **SRAM**: Static Random Access Memory
 - No refresh (6 transistors/bit vs. 1 transistor)
 - Size**: DRAM/SRAM - **4-8**,
 - Cost/Cycle time**: SRAM/DRAM - **8-16**

Main Memory Deep Background

- “Out-of-Core”, “In-Core,” “Core Dump”?
- “Core memory”?
- Non-volatile, magnetic
- Lost to 4 Kbit DRAM (today using 64Kbit DRAM)
- Access time 750 ns, cycle time 1500-3000 ns



DRAM logical organization (4 Mbit)



- Square root of bits per RAS/CAS

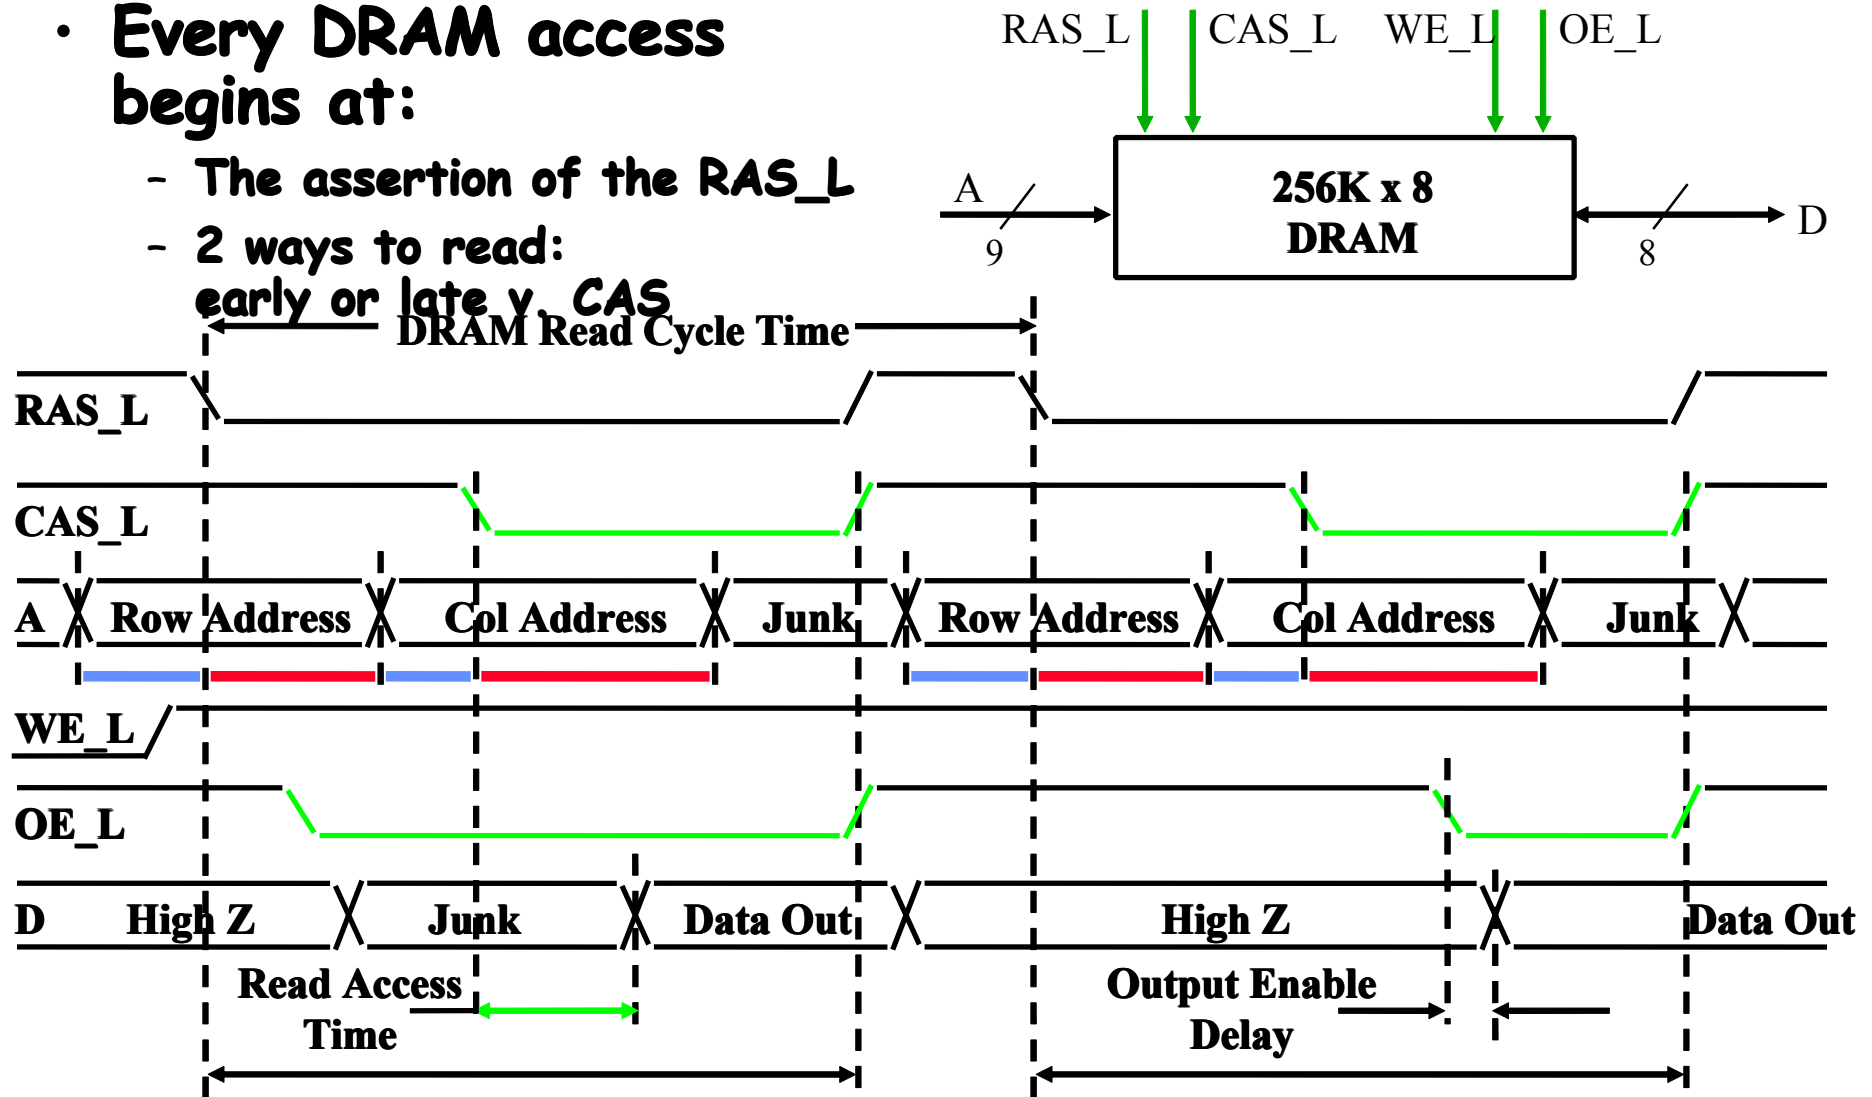
4 Key DRAM Timing Parameters

- **t_{RAC}** : minimum time from RAS line falling to the valid data output.
 - Quoted as the speed of a DRAM when buy
 - A typical 4Mb DRAM $t_{RAC} = 60$ ns
 - Speed of DRAM since on purchase sheet?
- **t_{RC}** : minimum time from the start of one row access to the start of the next.
 - $t_{RC} = 110$ ns for a 4Mbit DRAM with a t_{RAC} of 60 ns
- **t_{CAC}** : minimum time from CAS line falling to valid data output.
 - 15 ns for a 4Mbit DRAM with a t_{RAC} of 60 ns
- **t_{PC}** : minimum time from the start of one column access to the start of the next.
 - 35 ns for a 4Mbit DRAM with a t_{RAC} of 60 ns

DRAM Read Timing

- Every DRAM access begins at:

- The assertion of the RAS_L
- 2 ways to read:
early or late v. CAS



Early Read Cycle: OE_L asserted before CAS_L

Late Read Cycle: OE_L asserted after CAS_L

DRAM Performance

- A 60 ns (t_{RAC}) DRAM can
 - perform a row access only every 110 ns (t_{RC})
 - perform column access (t_{CAC}) in 15 ns, but time between column accesses is at least 35 ns (t_{PC}).
 - » In practice, external address delays and turning around buses make it 40 to 50 ns
- These times do not include the time to drive the addresses off the microprocessor nor the memory controller overhead!

DRAM History

- **DRAMs: capacity +60%/yr, cost -30%/yr**
 - 2.5X cells/area, 1.5X die size in -3 years
- **'98 DRAM fab line costs \$2B**
 - DRAM only: density, leakage v. speed
- **Rely on increasing no. of computers & memory per computer (60% market)**
 - SIMM or DIMM is replaceable unit
=> computers use any generation DRAM
- **Commodity, second source industry**
=> high volume, low profit, conservative
 - Little organization innovation in 20 years
- **Order of importance: 1) Cost/bit 2) Capacity**
 - First RAMBUS: 10X BW, +30% cost => little impact

DRAM Future: 1 Gbit DRAM (ISSCC '96; production '02?)


	Mitsubishi	Samsung
• Blocks	512 x 2 Mbit	1024 x 1 Mbit
• Clock	200 MHz	250 MHz
• Data Pins	64	16
• Die Size	24 x 24 mm	31 x 21 mm
- Sizes will be much smaller in production		
• Metal Layers	3	4
• Technology	0.15 micron	0.16 micron

Fast Memory Systems: DRAM specific

- Multiple CAS accesses: several names (page mode)
 - **Extended Data Out (EDO)**: 30% faster in page mode
- New DRAMs to address gap; what will they cost, will they survive?
 - **RAMBUS**: startup company; reinvent DRAM interface
 - » Each Chip a module vs. slice of memory
 - » Short bus between CPU and chips
 - » Does own refresh
 - » Variable amount of data returned
 - » 1 byte / 2 ns (500 MB/s per chip)
 - » 20% increase in DRAM area
 - **Synchronous DRAM**: 2 banks on chip, a clock signal to DRAM, transfer synchronous to system clock (66 - 150 MHz)
 - Intel claims Rambus Direct (16 b wide) is future PC memory?
 - » Possibly not true! Intel to drop Rambus?
- Niche memory or main memory?
 - e.g., Video RAM for frame buffers, DRAM + fast serial output

Main Memory Organizations

- **Simple:**
 - CPU, Cache, Bus, Memory same width (32 or 64 bits)
- **Wide:**
 - CPU/Mux 1 word; Mux/Cache, Bus, Memory N words (Alpha: 64 bits & 256 bits; UltraSPARC 512)
- **Interleaved:**
 - CPU, Cache, Bus 1 word; Memory N Modules (4 Modules); example is *word interleaved*





Main Memory Performance

- Timing model (word size is 32 bits)
 - 1 to send address,
 - 6 access time, 1 to send data
 - Cache Block is 4 words
- **Simple M.P.** $= 4 \times (1+6+1) = 32$
- **Wide M.P.** $= 1 + 6 + 1 = 8$
- **Interleaved M.P.** $= 1 + 6 + 4 \times 1 = 11$

Address++	Bank 0	Address++	Bank 1	Address++	Bank 2	Address++	Bank 3
0		1		2		3	
4		5		6		7	
8		9		10		11	
12		13		14		15	

Independent Memory Banks

- Memory banks for independent accesses vs. faster sequential accesses
 - Multiprocessor
 - I/O
 - CPU with Hit under n Misses, Non-blocking Cache
- **Superbank**: all memory active on one block transfer (or **Bank**)
- **Bank**: portion within a superbank that is word interleaved (or **Subbank**)



Independent Memory Banks

- **How many banks?**
number banks \leq number clocks to access word in bank
 - For sequential accesses, otherwise will return to original bank before it has next word ready
 - (like in vector case)
- **Increasing DRAM \Rightarrow fewer chips \Rightarrow harder to have banks**

Avoiding Bank Conflicts

- Lots of banks

```
int x[256][512];  
    for (j = 0; j < 512; j = j+1)  
        for (i = 0; i < 256; i = i+1)  
            x[i][j] = 2 * x[i][j];
```

- Even with 128 banks, since 512 is multiple of 128, conflict on word accesses
- SW: loop interchange or declaring array not power of 2 ("array padding")
- HW: Prime number of banks
 - bank number = address mod number of banks
 - address within bank = address / number of words in bank
 - modulo & divide per memory access with prime no. banks?
 - address within bank = address mod number words in bank
 - bank number? easy if 2^N words per bank

Fast Bank Number

- Chinese Remainder Theorem

As long as two sets of integers a_i and b_i follow these rules

$$b_i = x \bmod a_i, 0 \leq b_i < a_i, 0 \leq x < a_0 \times a_1 \times a_2 \times \dots$$

and that a_i and a_j are co-prime if $i \neq j$, then the integer x has only one solution (unambiguous mapping):

- bank number = b_0 , number of banks = a_0 (= 3 in example)
- address within bank = b_1 , number of words in bank = a_1 (= 8 in example)
- N word address 0 to N-1, prime no. banks, words power of 2

		Seq. Interleaved			Modulo Interleaved		
Bank Number:		0	1	2	0	1	2
Address							
within Bank:							
0	0	0	1	2	0	16	8
1	3	4	5	6	9	1	17
2	6	7	8	9	18	10	2
3	9	10	11	12	3	19	11
4	12	13	14	15	12	4	20
5	15	16	17	18	21	13	5
6	18	19	20	21	6	22	14
7	21	22	23	24	15	7	23

DRAMs per PC over Time

		DRAM Generation					
		'86	'89	'92	'96	'99	'02
		1 Mb	4 Mb	16 Mb	64 Mb	256 Mb	1 Gb
Minimum Memory Size	4 MB	32 → 8					
	8 MB		16 → 4				
	16 MB			8 → 2			
	32 MB				4 → 1		
	64 MB				8 → 2		
	128 MB					4 → 1	
	256 MB					8 → 2	

Need for Error Correction!

- **Motivation:**
 - Failures/time *proportional* to number of bits!
 - As DRAM cells shrink, more vulnerable
- **Went through period in which failure rate was low enough without error correction that people didn't do correction**
 - DRAM banks too large now
 - Servers always corrected memory systems
- **Basic idea: add redundancy through parity bits**
 - Simple but wasteful version:
 - » Keep three copies of everything, vote to find right value
 - » 200% overhead, so not good!
 - Common configuration: Random error correction
 - » SEC-DED (single error correct, double error detect)
 - » One example: 64 data bits + 8 parity bits (11% overhead)
 - » Papers up on reading list from last term tell you how to do these types of codes
 - Really want to handle failures of physical components as well
 - » Organization is multiple DRAMs/SIMM, multiple SIMMs
 - » Want to recover from failed DRAM and failed SIMM!
 - » Requires more redundancy to do this
 - » All major vendors thinking about this in high-end machines

Architecture in practice

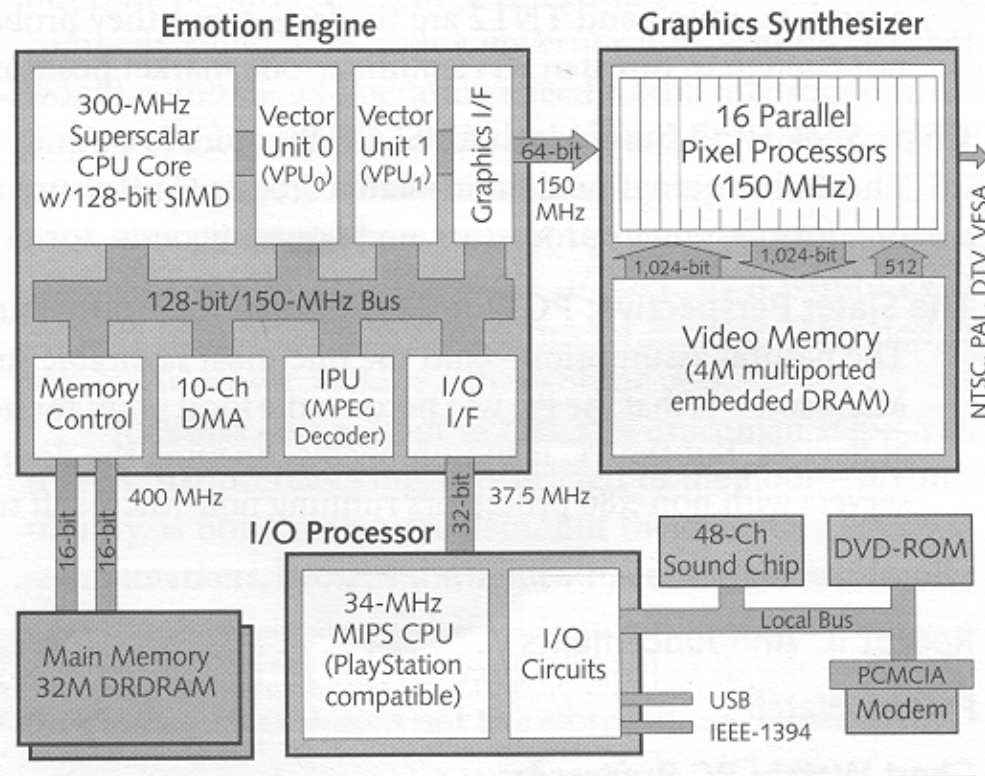


Figure 1. PlayStation 2000 employs an unprecedented level of parallelism to achieve workstation-class 3D performance.




Figure 2. PlayStation 2000 screenshot. (Source: Namco)

- (as reported in **Microprocessor Report**, Vol 13, No. 5)
 - **Emotion Engine: 6.2 GFLOPS, 75 million polygons per second**
 - **Graphics Synthesizer: 2.4 Billion pixels per second**
 - **Claim: *Toy Story* realism brought to games!**

More esoteric Storage Technologies?


- **Tunneling Magnetic Junction RAM (TMJ-RAM):**
 - Speed of SRAM, density of DRAM, non-volatile (no refresh)
 - New field called "Spintronics": combination of quantum spin and electronics
 - Same technology used in high-density disk-drives
- **MEMs storage devices:**
 - Large magnetic "sled" floating on top of lots of little read/write heads
 - Micromechanical actuators move the sled back and forth over the heads

Tunneling Magnetic Junction



MEMS-based Storage

- **Magnetic “sled” floats on array of read/write heads**
 - Approx 250 Gbit/in²
 - Data rates:
IBM: 250 MB/s w 1000 heads
CMU: 3.1 MB/s w 400 heads
- **Electrostatic actuators move media around to align it with heads**
 - Sweep sled $\pm 50\mu\text{m}$ in $< 0.5\mu\text{s}$
- **Capacity estimated to be in the 1-10GB in 10cm²**



See Ganger et al: <http://www.lcs.ece.cmu.edu/research/MEMS>

Main Memory Summary

- **Wider Memory**
- **Interleaved Memory: for sequential or independent accesses**
- **Avoiding bank conflicts: SW & HW**
- **DRAM specific optimizations: page mode & Specialty DRAM**
- **Need Error correction**

Review: Improving Cache Performance


1. Reduce the miss rate,
2. Reduce the miss penalty, or
3. Reduce the time to hit in the cache.

$$AMAT = \text{HitTime} + MissRate \times MissPenalty$$


1. Fast Hit times via Small and Simple Caches

- **Why Alpha 21164 has 8KB Instruction and 8KB data cache + 96KB second level cache?**
 - Small data cache and clock rate
- **Direct Mapped, on chip**


2. Fast hits by Avoiding Address Translation



Conventional Organization



Virtually Addressed Cache
Translate only on miss
Synonym Problem




Overlap \$ access
with VA translation:
requires \$ index to
remain invariant
across translation

2. Fast hits by Avoiding Address Translation

- Send virtual address to cache? Called Virtually Addressed Cache or just Virtual Cache vs. Physical Cache
 - Every time process is switched logically must flush the cache; otherwise get false hits
 - » Cost is time to flush + “compulsory” misses from empty cache
 - Dealing with aliases (sometimes called synonyms):
Two different virtual addresses map to same physical address
 - I/O must interact with cache, so need virtual address
- Solution to aliases
 - HW guarantees covers index field & direct mapped, they must be unique;
called page coloring
- Solution to cache flush
 - Add process identifier tag that identifies process as well as address within process: can't get a hit if wrong process

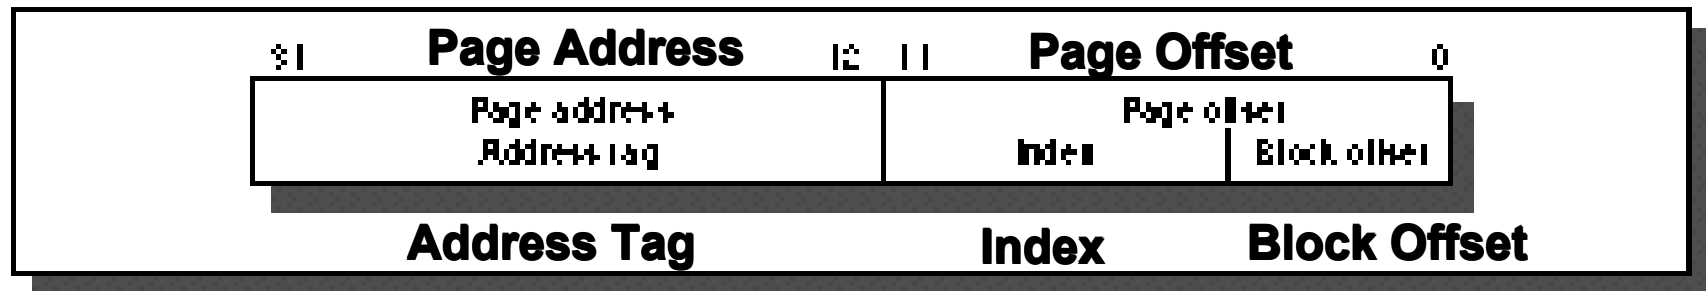
2. Fast Cache Hits by Avoiding Translation: Process ID impact

- Black is uniprocess
- Light Gray is multiprocess when flush cache
- Dark Gray is multiprocess when use Process ID tag
- Y axis: Miss Rates up to 20%
- X axis: Cache size from 2 KB to 1024 KB



2. Fast Cache Hits by Avoiding Translation: Index with Physical Portion of Address

- If index is physical part of address, can start tag access in parallel with translation so that can compare to physical tag



- Limits cache to page size: what if want bigger caches and uses same trick?
 - Higher associativity moves barrier to right
 - Page coloring

3: Fast Hits by pipelining Cache

Case Study: MIPS R4000

- **8 Stage Pipeline:**
 - IF-first half of fetching of instruction; PC selection happens here as well as initiation of instruction cache access.
 - IS-second half of access to instruction cache.
 - RF-instruction decode and register fetch, hazard checking and also instruction cache hit detection.
 - EX-execution, which includes effective address calculation, ALU operation, and branch target computation and condition evaluation.
 - DF-data fetch, first half of access to data cache.
 - DS-second half of access to data cache.
 - TC-tag check, determine whether the data cache access hit.
 - WB-write back for loads and register-register operations.
- **What is impact on Load delay?**
 - Need 2 instructions between a load and its use!

Case Study: MIPS R4000

**TWO Cycle
Load Latency**

IF	IS	RF	EX	DF	DS	TC	WB
	IF	IS	RF	EX	DF	DS	TC
		IF	IS	RF	EX	DF	DS
			IF	IS	RF	EX	DF
				IF	IS	RF	EX
					IF	IS	RF
						IF	IS
							IF

**THREE Cycle
Branch Latency**

(conditions evaluated
during EX phase)

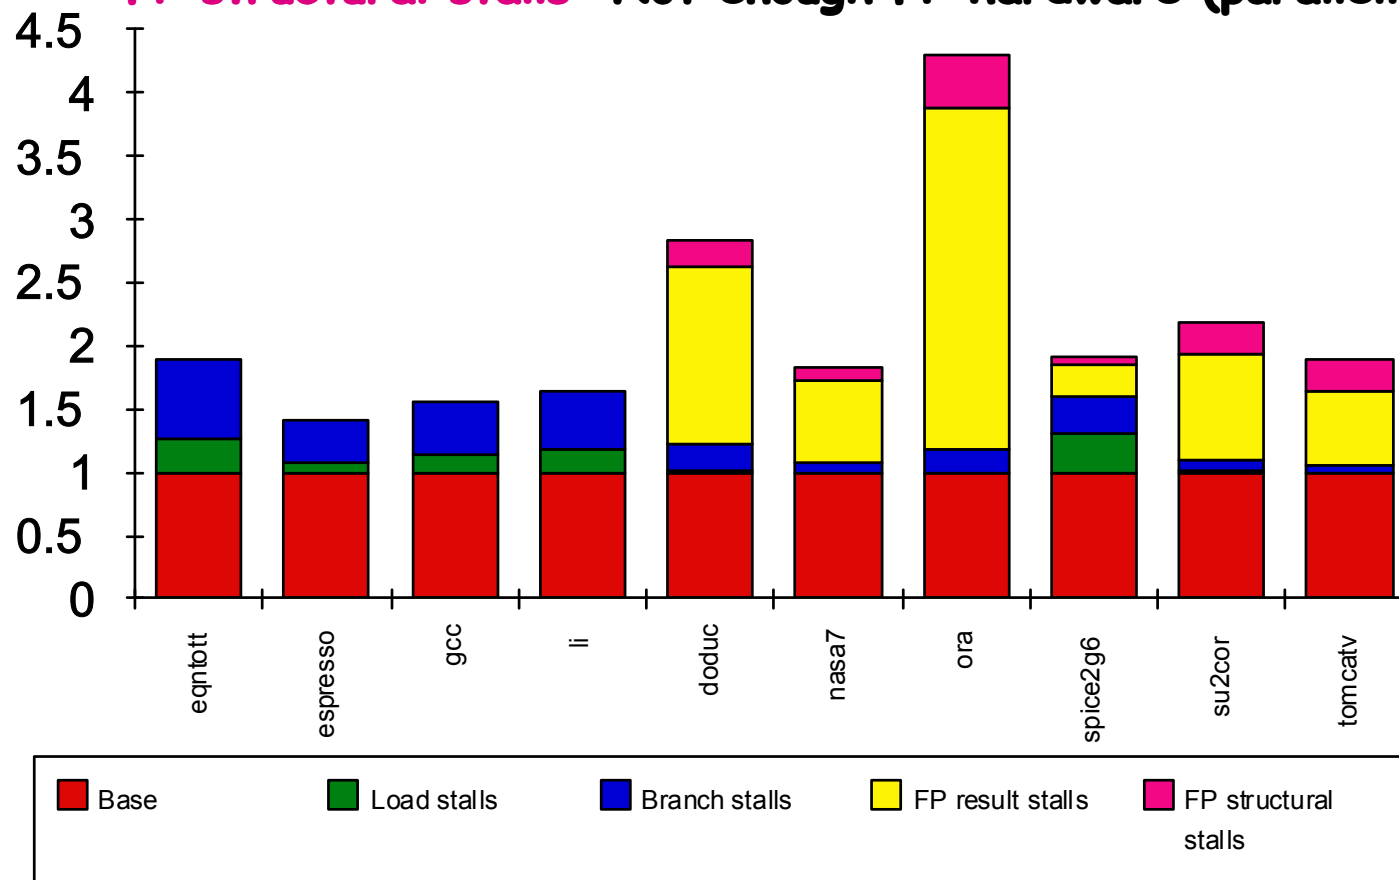
Delay slot plus two stalls

Branch likely cancels delay slot if not taken

IF	IS	RF	EX	DF	DS	TC	WB
	IF	IS	RF	EX	DF	DS	TC
		IF	IS	RF	EX	DF	DS
			IF	IS	RF	EX	DF
				IF	IS	RF	EX
					IF	IS	RF
						IF	IS
							IF


R4000 Performance

- Not ideal CPI of 1:
 - Load stalls (1 or 2 clock cycles)
 - Branch stalls (2 cycles + unfilled slots)
 - FP result stalls: RAW data hazard (latency)
 - FP structural stalls: Not enough FP hardware (parallelism)



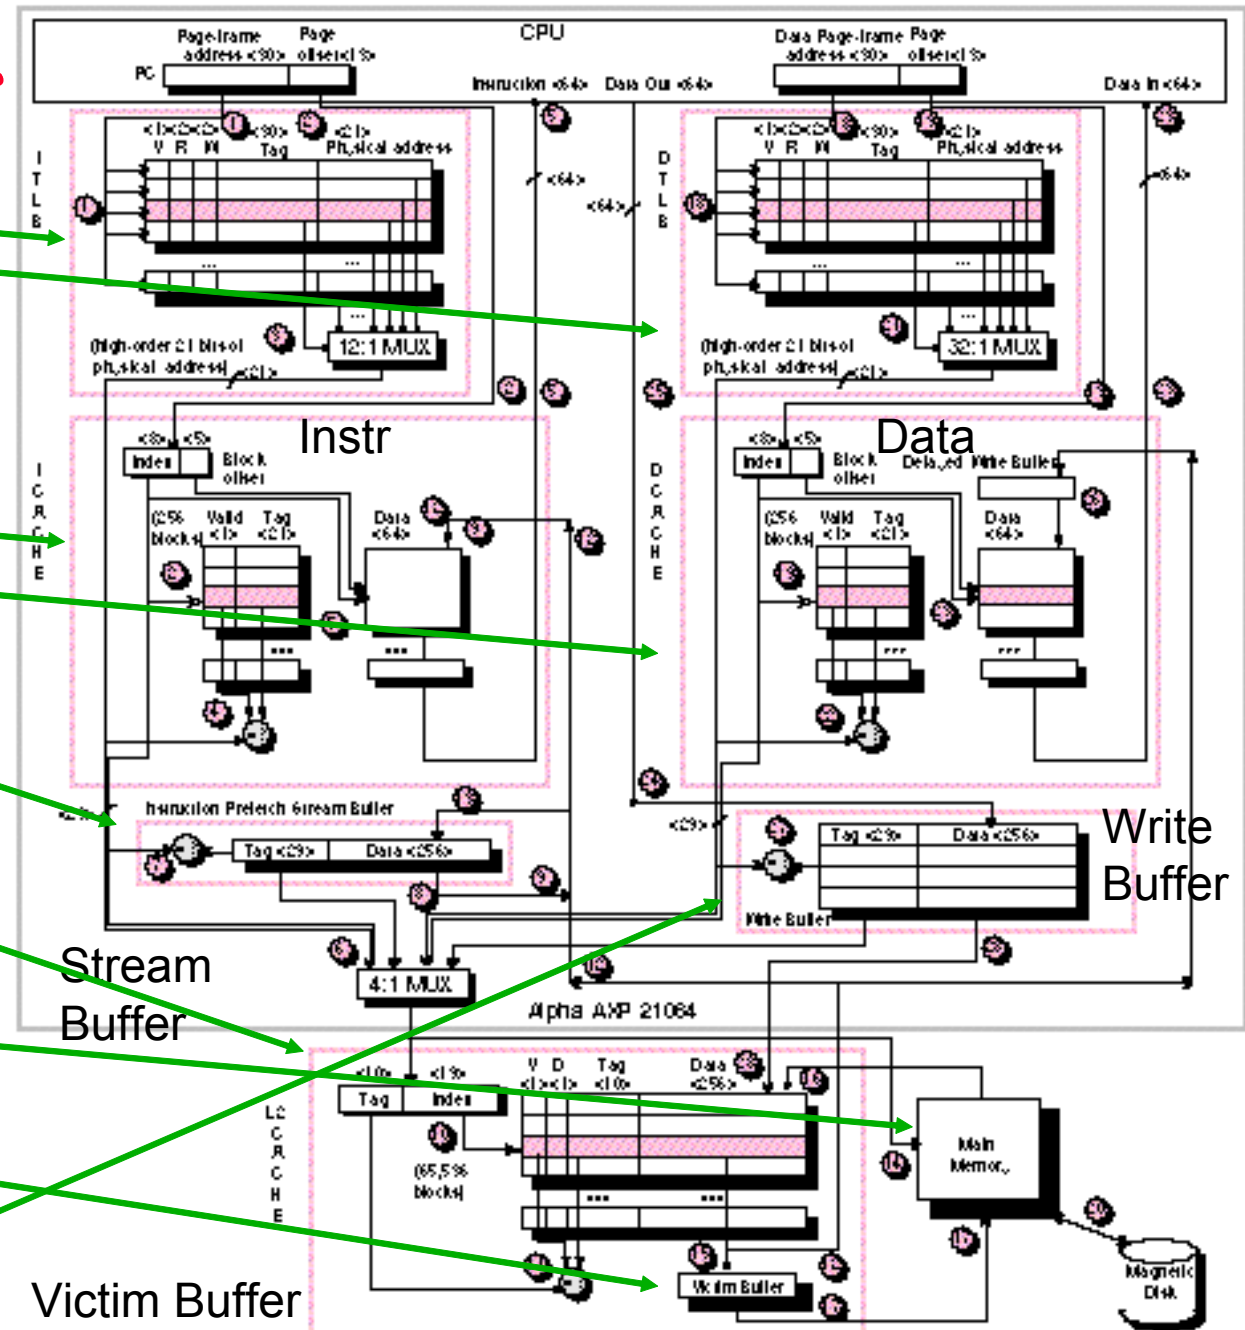
What is the Impact of What You've Learned About Caches?

- 1960-1985: Speed = $f(\text{no. operations})$
- 1990
 - Pipelined Execution & Fast Clock Rate
 - Out-of-Order execution
 - Superscalar Instruction Issue
- 1998: Speed = $f(\text{non-cached memory accesses})$
- What does this mean for
 - Compilers?, Operating Systems?, Algorithms? Data Structures?

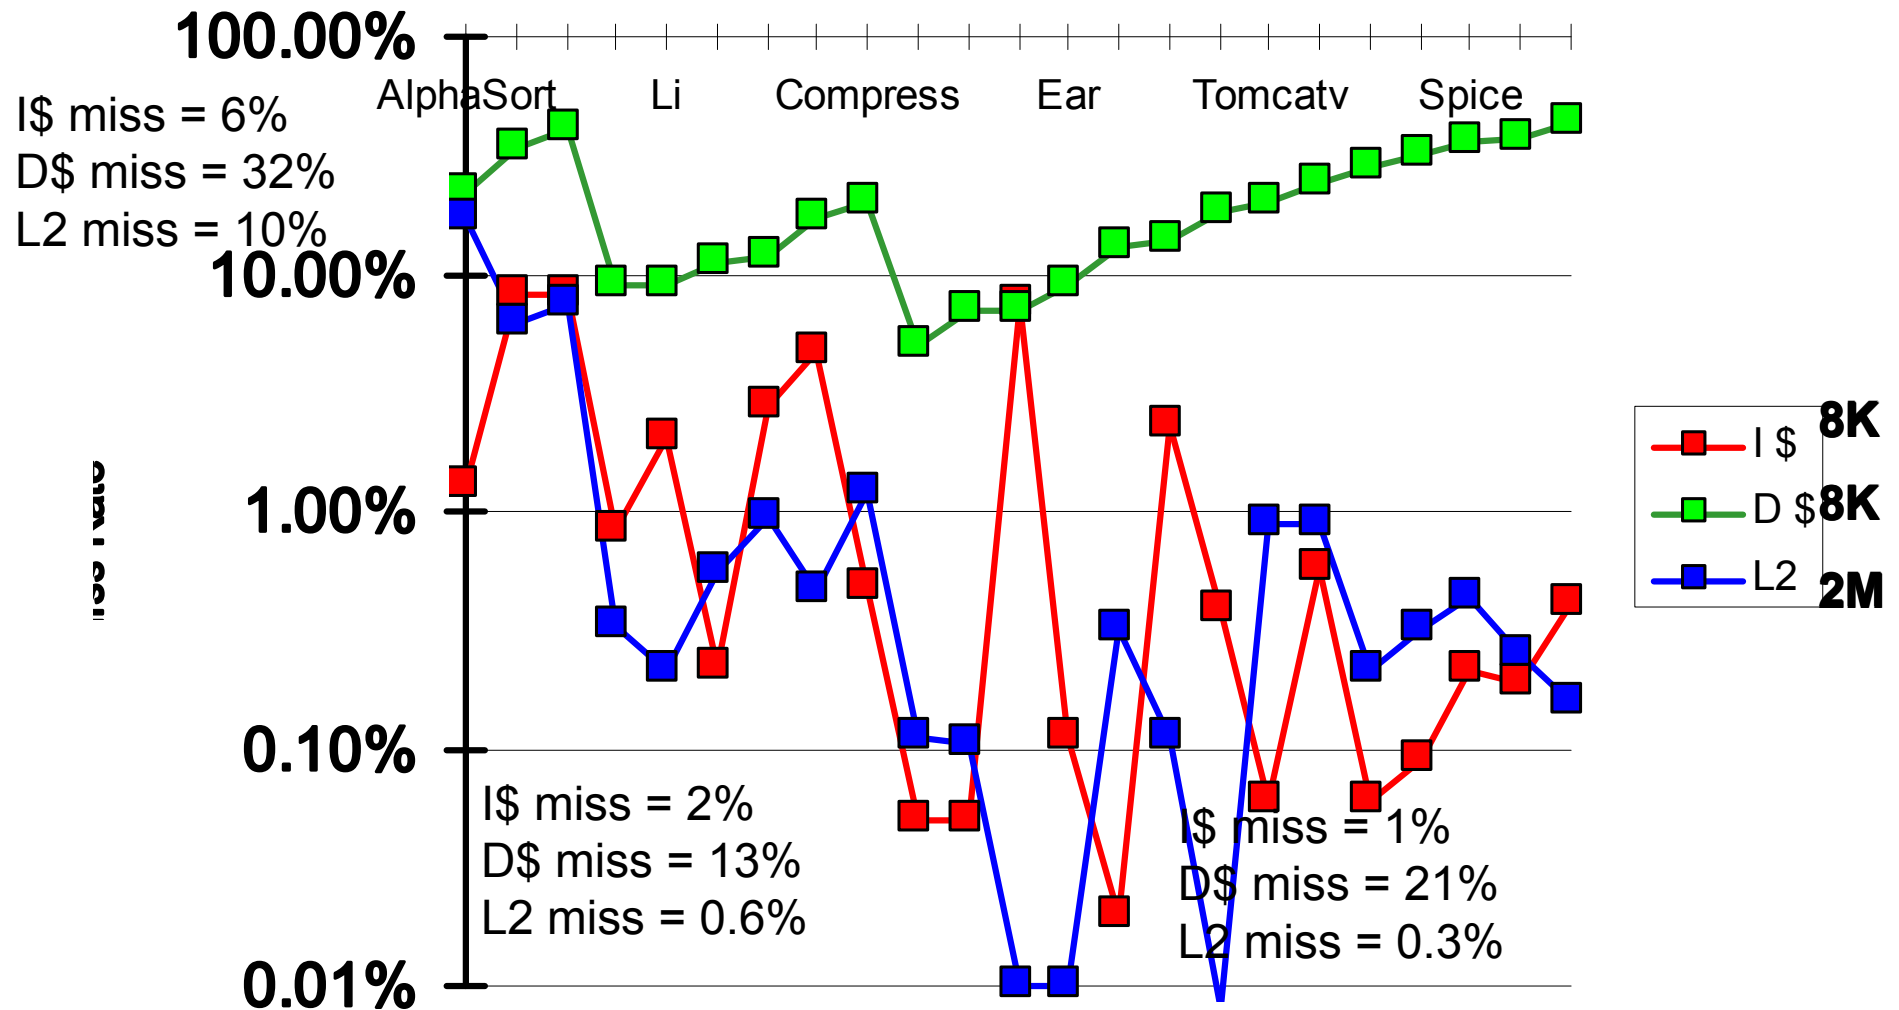


Alpha 21064

- Separate Instr & Data TLB & Caches
- TLBs fully associative
- TLB updates in SW ("Priv Arch Lib")
- Caches 8KB direct mapped, write thru
- Critical 8 bytes first
- Prefetch instr. stream buffer
- 2 MB L2 cache, direct mapped, WB (off-chip)
- 256 bit path to main memory, 4 x 64-bit modules
- Victim Buffer: to give read priority over write
- 4 entry write buffer between D\$ & L2\$

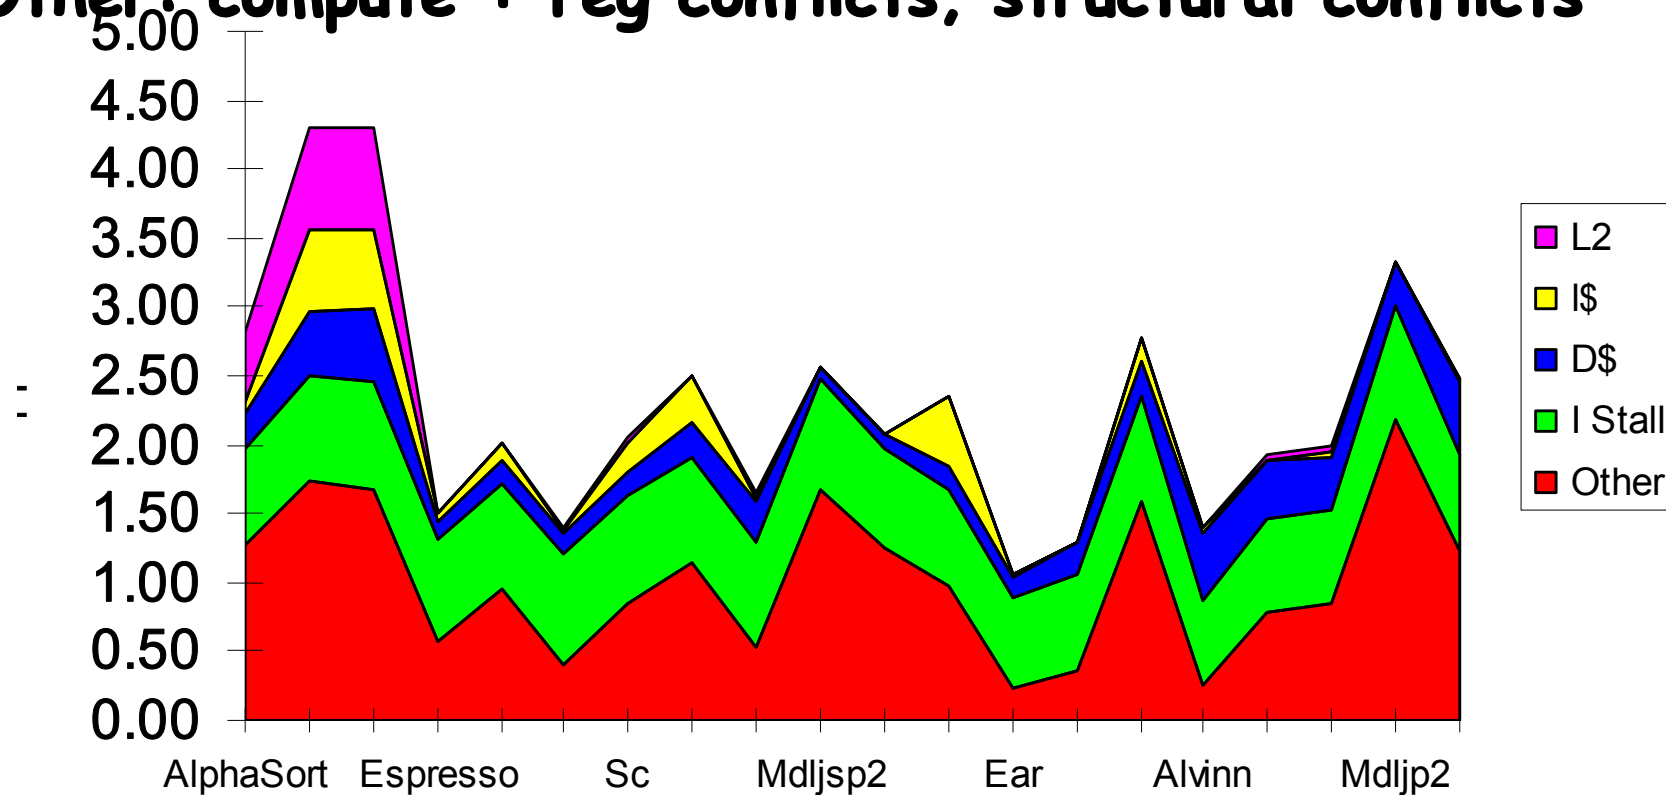


Alpha Memory Performance: Miss Rates of SPEC92



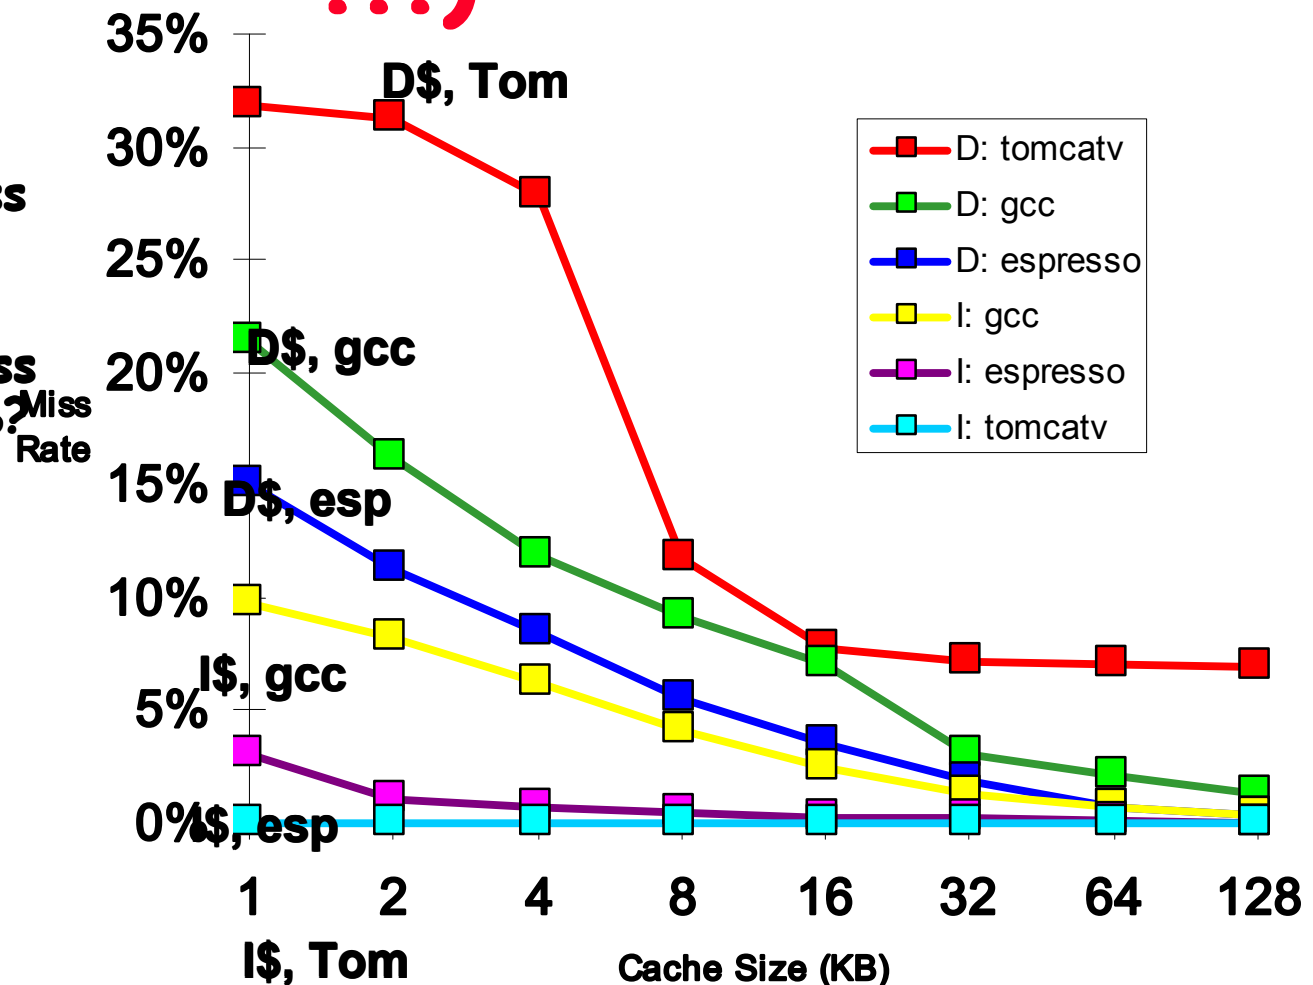
Alpha CPI Components

- Instruction stall: branch mispredict (green);
- Data cache (blue); Instruction cache (yellow); L2\$ (pink)
- Other: compute + reg conflicts, structural conflicts



Pitfall: Predicting Cache Performance from Different Prog. (ISA, compiler,)

- 4KB Data cache miss rate 8%, 12%, or 28%?
- 1KB Instr cache miss rate 0%, 3%, or 10%?
- Alpha vs. MIPS for 8KB Data \$: 17% vs. 10%
- Why 2X Alpha v. MIPS?



Cache Optimization Summary

	Technique	MR	MP	HT	Complexity
miss rate	Larger Block Size	+	-		0
	Higher Associativity	+		-	1
	Victim Caches	+			2
	Pseudo-Associative Caches	+			2
	HW Prefetching of Instr/Data	+			2
	Compiler Controlled Prefetching	+			3
	Compiler Reduce Misses	+			0
miss penalty	Priority to Read Misses		+		1
	Early Restart & Critical Word 1st		+		2
	Non-Blocking Caches		+		3
	Second Level Caches		+		2
	Better memory system		+		3
hit time	Small & Simple Caches	-		+	0
	Avoiding Address Translation			+	2
	Pipelining Caches			+	2